

```
In [ ]: import pandas as pd
pd.set_option("display.max_columns", None)

import matplotlib.pyplot as plt

Model Building
```

```
In [ ]: df = pd.read_csv("dataset/phase2_df.csv")

# Replace spaces with underscores in column names
df.rename(columns=lambda x: x.replace(' ', '_'), inplace=True)

df.head()
```

Out []:

	Age	Gender	Education_Level	Joining_Designation	joining_date	joining_month	joining_year	ReportCount	Active	TBV_avg	QR_avg	Income_avg	Grade_avg	income_diff_pattern	tbv_diff_pattern	qr_diff_pattern	grade_diff_pattern	City_C1	City_C10	City_C11	City_C12	City_C13	City_C14
0	0.189189	0.0	1.0	0.00	0.766667	1.000000	0.714286	0.086957	0.0	0.184600	0.333333	0.262508	0.00	0.0	0.000000	0.500000	0.0	0	0	0	0	0	0
1	0.216216	0.0	0.0	0.00	0.000000	0.727273	0.857143	0.086957	0.0	0.057086	0.000000	0.200489	0.00	0.0	0.333333	0.500000	0.0	0	0	0	0	0	0
2	0.232667	0.0	1.0	0.00	0.900000	0.363636	0.285714	0.956522	0.0	0.153949	0.086957	0.610567	0.75	0.0	0.333333	0.500000	0.0	0	0	0	0	0	0
3	0.243243	1.0	0.0	0.25	0.966667	0.909091	0.714286	0.043478	0.0	0.089047	0.000000	0.237608	0.25	0.0	0.333333	0.500000	0.0	0	0	0	0	0	0
4	0.571429	0.0	1.0	0.00	0.066667	0.454545	0.714286	0.260870	0.0	0.082327	0.142857	0.227116	0.00	0.0	0.333333	0.333333	0.0	0	0	0	0	0	0

```
In [ ]: df.shape
```

Out []: (2381, 46)

```
In [ ]: from sklearn.model_selection import train_test_split

# Separate features (X) and target variable (y)
X = df.drop(columns=['Active'])
y = df['Active'] # Fully-Paid: 1 and Charged-Off: 0
# Perform train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

```
In [ ]: y_train.value_counts(), y_test.value_counts()
```

Out []:

```
(Active
0.0    1440
1.0     702
Name: count, dtype: int64,
Active
0.0     176
1.0      63
Name: count, dtype: int64)
```

- It doesn't seem to be a big data imbalance so using techniques like SMOTE or Weight Imbalance will have a big impact.
- Also using such techniques may increase the chances of overfitting.

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Initialize RandomForestClassifier
clf = RandomForestClassifier(random_state=42)

# Fit the model
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:\n", report)
```

Accuracy: 0.91

Confusion Matrix:

```
[[166  10]
 [ 12  51]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.93	0.94	0.94	176
1.0	0.84	0.81	0.82	63
accuracy			0.91	239
macro avg	0.88	0.88	0.88	239
weighted avg	0.91	0.91	0.91	239

Let's use Class weights to balance the Dataset and check the results

```
In [ ]: # Specify class weights (class 0 is the minority class)
class_weights = {0: 1, 1: 2}

# Initialize RandomForestClassifier
clf = RandomForestClassifier(class_weight=class_weights, random_state=42)

# Fit the model
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:\n", report)
```

Accuracy: 0.92

Confusion Matrix:

```
[[169   7]
 [ 12  51]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.93	0.96	0.95	176
1.0	0.88	0.81	0.84	63
accuracy			0.92	239
macro avg	0.91	0.88	0.89	239
weighted avg	0.92	0.92	0.92	239

- 1% increase in accuracy using Class Weights

```
In [ ]: from imblearn.over_sampling import SMOTE

# Apply SMOTE to oversample the minority class
smote = SMOTE(sampling_strategy='auto', random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
# Check the distribution of classes after applying SMOTE
print("Class distribution before SMOTE:", {0: sum(y_train == 0), 1: sum(y_train == 1)})
print("Class distribution after SMOTE:", {0: sum(y_train_smote == 0), 1: sum(y_train_smote == 1)})

Class distribution before SMOTE: {0: 1440, 1: 702}
Class distribution after SMOTE: {0: 1440, 1: 1440}
```

```
In [ ]: # Initialize RandomForestClassifier
clf = RandomForestClassifier(random_state=42)

# Fit the model
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:\n", report)
```

Accuracy: 0.91

Confusion Matrix:
[[166 10]
 [12 51]]

Classification Report:				
	precision	recall	f1-score	support
0.0	0.93	0.94	0.94	176
1.0	0.84	0.81	0.82	63
accuracy			0.91	239
macro avg	0.88	0.88	0.88	239
weighted avg	0.91	0.91	0.91	239

- Not much change if we use SMOTE

Hyperparameter Tuning

```
In [ ]: from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

# Define the parameter grid
param_dist = {
    'n_estimators': randint(100, 300),
    'max_depth': [None] + list(range(25, 40, 1)),
    'min_samples_split': [5,10,20,25],
    'min_samples_leaf': [1, 2, 3],
    'bootstrap': [False]
}

# Initialize RandomForestClassifier
rf_clf = RandomForestClassifier(class_weight=class_weights, random_state=42)

# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(
    rf_clf, param_distributions=param_dist, n_iter=10, cv=5, scoring='accuracy', random_state=42, n_jobs=-1
)

# Fit the model
random_search.fit(X_train, y_train)

# Display the best hyperparameters
print("Best Hyperparameters:", random_search.best_params_)

# Make predictions on the test set using the best model
y_pred = random_search.best_estimator_.predict(X_test)

# Print confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:\n", pd.DataFrame(conf_matrix, columns=['Predicted 0', 'Predicted 1'], index=['Actual 0', 'Actual 1']))

# Print accuracy and classification report
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"\nAccuracy: {accuracy:.2f}")
print("\nClassification Report:\n", report)
```

Best Hyperparameters: {'bootstrap': False, 'max_depth': 39, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 157}

Confusion Matrix:
Predicted 0 Predicted 1
Actual 0 167 9
Actual 1 11 52

Accuracy: 0.92

Classification Report:				
	precision	recall	f1-score	support
0.0	0.94	0.95	0.94	176
1.0	0.85	0.83	0.84	63
accuracy			0.92	239
macro avg	0.90	0.89	0.89	239
weighted avg	0.92	0.92	0.92	239

- Accuracy 91: Best Hyperparameters: {'bootstrap': False, 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 68}
- Accuracy 91: Best Hyperparameters: {'bootstrap': False, 'max_depth': 23, 'min_samples_leaf': 2, 'min_samples_split': 20, 'n_estimators': 197}
- Accuracy 92: Best Hyperparameters: {'bootstrap': False, 'max_depth': 26, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 112}
- Accuracy 92: Best Hyperparameters: {'bootstrap': False, 'max_depth': 39, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 157}

Implementation of Boosting Algorithms

```
In [ ]: from sklearn.ensemble import GradientBoostingClassifier

# Specify class weights (class 0 is the minority class)
class_weights = {0: 1, 1: 2}

# Initialize RandomForestClassifier
clf = GradientBoostingClassifier(random_state=42)

# Fit the model
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:\n", report)
```

Accuracy: 0.93

Confusion Matrix:
[[168 8]
 [8 55]]

Classification Report:				
	precision	recall	f1-score	support
0.0	0.95	0.95	0.95	176
1.0	0.87	0.87	0.87	63
accuracy			0.93	239
macro avg	0.91	0.91	0.91	239
weighted avg	0.93	0.93	0.93	239

XGBoost

```
In [ ]: import xgboost as xgb

# Initialize XGBoost classifier
clf = xgb.XGBClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)

# Fit the model
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:\n", report)
```

Accuracy: 0.93

Confusion Matrix:
[[168 8]
 [8 55]]

Classification Report:				
	precision	recall	f1-score	support
0.0	0.95	0.95	0.95	176
1.0	0.87	0.87	0.87	63
accuracy			0.93	239
macro avg	0.91	0.91	0.91	239
weighted avg	0.93	0.93	0.93	239

```
In [ ]: ## Hyperparameter Tuning
# Define the parameter grid for XGBoost
param_dist = {
    'learning_rate': [0.008, 0.009, 0.01, 0.02],
    'n_estimators': [250, 260, 270, 280, 290],
    'max_depth': [6, 7, 8, 9],
    'subsample': [0.7, 0.8, 0.9, 1.0],
    'colsample_bytree': [0.8, 0.9, 1.0]
}
```

```
# Initialize XGBoost classifier
xgb_clf = xgb.XGBClassifier(random_state=42)

# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(
    xgb_clf, param_distributions=param_dist, n_iter=10, cv=5, scoring='accuracy', random_state=42, n_jobs=-1
)

# Fit the model
random_search.fit(X_train, y_train)

# Display the best hyperparameters
print("Best Hyperparameters:", random_search.best_params_)

# Make predictions on the test set using the best model
y_pred = random_search.best_estimator_.predict(X_test)

# Print confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:\n", pd.DataFrame(conf_matrix, columns=['Predicted 0', 'Predicted 1'], index=['Actual 0', 'Actual 1']))

# Print accuracy and classification report
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"\nAccuracy: {accuracy:.2f}")
print("\nClassification Report:\n", report)
```

Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 280, 'max_depth': 6, 'learning_rate': 0.02, 'colsample_bytree': 1.0}

Confusion Matrix:

	Predicted 0	Predicted 1
Actual 0	168	8
Actual 1	8	55

Accuracy: 0.93

Classification Report:

	precision	recall	f1-score	support
0.0	0.95	0.95	0.95	176
1.0	0.87	0.87	0.87	63
accuracy			0.93	239
macro avg	0.91	0.91	0.91	239
weighted avg	0.93	0.93	0.93	239

- Accuracy 92: Best Hyperparameters: {'subsample': 0.7, 'n_estimators': 200, 'max_depth': 9, 'learning_rate': 0.01, 'colsample_bytree': 1.0}
- Accuracy 93: Best Hyperparameters: {'subsample': 0.9, 'n_estimators': 250, 'max_depth': 7, 'learning_rate': 0.01, 'colsample_bytree': 1.0}
- Accuracy 93: Best Hyperparameters: {'subsample': 0.9, 'n_estimators': 250, 'max_depth': 8, 'learning_rate': 0.01, 'colsample_bytree': 0.8}
- Accuracy 93: Best Hyperparameters: {'subsample': 0.9, 'n_estimators': 280, 'max_depth': 6, 'learning_rate': 0.01, 'colsample_bytree': 0.9}

```
In [ ]: import xgboost as xgb

# Initialize XGBoost classifier
clf = xgb.XGBClassifier(subsample= 0.9, n_estimators= 280, max_depth= 6, learning_rate= 0.01, colsample_bytree= 0.9, random_state=42)

# Fit the model
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:\n", report)
```

Accuracy: 0.93

Confusion Matrix:

[[169	7]
[10	53]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.94	0.96	0.95	176
1.0	0.88	0.84	0.86	63
accuracy			0.93	239
macro avg	0.91	0.90	0.91	239
weighted avg	0.93	0.93	0.93	239

LightGBM

```
In [ ]: import lightgbm as lgb

# Initialize LightGBM classifier
clf = lgb.LGBMClassifier(random_state=42)

# Fit the model
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"\nAccuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:\n", report)
```

[LightGBM] [Info] Number of positive: 702, number of negative: 1440
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000236 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 1022
[LightGBM] [Info] Number of data points in the train set: 2142, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.327731 -> initscore=-0.718465
[LightGBM] [Info] Start training from score -0.718465

Accuracy: 0.94

Confusion Matrix:

[[168	8]
[7	56]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.96	0.95	0.96	176
1.0	0.88	0.89	0.88	63
accuracy			0.94	239
macro avg	0.92	0.92	0.92	239
weighted avg	0.94	0.94	0.94	239


```
In [ ]: from sklearn.metrics import roc_curve, auc

# Get predicted probabilities for the positive class (class 1)
y_prob = clf.predict_proba(X_test)[:, 1]

# Compute ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)

# Compute Area Under the Curve (AUC)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```

Feature Importance

- Since LightGBM is giving us the best accuracy, we'll be using the LightGBM trained model to extract feature importance.

```
In [ ]: # Extract feature importance
feature_importance = clf.feature_importances_

# Create a DataFrame to display feature names and their importance scores
feature_importance_df = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': feature_importance
})

# Sort the DataFrame by importance score in descending order
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)

# Display the feature importance DataFrame
print(feature_importance_df)
```

	Feature	Importance
7	ReportCount	451
4	joining_date	404
0	Age	391
10	Income_avg	339
5	joining_month	298
6	joining_year	233
8	TBV_avg	215
13	tbv_diff_pattern	120
9	QR_avg	117
14	qr_diff_pattern	67
2	Education_Level	55
1	Gender	49
3	Joining_Designation	44
35	City_C27	24
11	Grade_avg	24
37	City_C29	15
18	City_C11	13
40	City_C5	13
29	City_C21	11
24	City_C17	11
16	City_C1	11
34	City_C26	11
20	City_C13	9
26	City_C19	8
25	City_C18	8
36	City_C28	7
38	City_C3	7
30	City_C22	6
42	City_C7	5
27	City_C2	5
41	City_C6	5
19	City_C12	4
31	City_C23	4
32	City_C24	4
33	City_C25	3
28	City_C20	3
17	City_C10	2
44	City_C9	2
43	City_C8	1
23	City_C16	1
12	income_diff_pattern	0
22	City_C15	0
21	City_C14	0
15	grade_diff_pattern	0
39	City_C4	0

Results Interpretation & Stakeholder Presentation

Interpreting Model Coefficients: It seems that top features influencing the usecase output includes:

- ReportCount: Number of time the Driver is reporting.
- Joining Date: It looks interesting but Joining Date has good influence on output.
- Age: Driver's age across the Business has good influence on the output
- Income_avg: What income the Driver is earning from the Business.
- Joining_month: Interestingly joining month also has an impact on Driver's attrition.

Feedback Loop

- We should set up a periodic review process to assess the model's relevance and performance.
- Keep on generating more data and analyse the reports on what may impacting out Churing behaviour and handle the same with care and understanding.
- Recommend surveys or feedback mechanisms to collect data on new trends, driver concerns, and customer feedback to refine the model in the future.
- We should have some daily feedback calls form driver, understanding their problems and saving their respective sentiments.

Trade-Off Analysis and Recommendation is in Notebook Phase 2

```
In [ ]:
```